

# Decision Theory

## Lecture 10

# Time to make a decision...

Exercise inspired by Mausam, University of Washington, CSE573

## State of Nature

Poor market performance    Good market performance

**Payoff**

**Payoff**

Buy Apple	-1,000	1,700
Buy Google	-2,000	2,100
Don't invest	10	10

# How to invest \$10 in 2004?

# Maximax

Optimism

Select the maximum of the maximum payoff

Action

	State of Nature		Criterion
	Poor market performance	Good market performance	Maximum payoff for an action
	Payoff	Payoff	
Buy Apple	-1,000	1,700	1,700
Buy Google	-2,000	2,100	2,100
Don't invest	10	10	10

# Maximax

Optimism

Select the maximum of the maximum payoff

Action

	State of Nature		Criterion
	Poor market performance Payoff	Good market performance Payoff	Maximum payoff for an action
Buy Apple	-1,000	1,700	1,700
Buy Google	-2,000	2,100	2,100
Don't invest	10	10	10

← **Maximax**

# Maximin

Pessimism

Select the maximum of the minimum payoffs

Action

	State of Nature		Criterion
	Poor market performance Payoff	Good market performance Payoff	Minimum payoff for an action
Buy Apple	-1,000	1,700	-1,000
Buy Google	-2,000	2,100	-2,000
Don't invest	10	10	10

# Maximin

Pessimism

Select the maximum of the minimum payoffs

Action

	State of Nature		Criterion
	Poor market performance	Good market performance	Minimum payoff for an action
	Payoff	Payoff	
Buy Apple	-1,000	1,700	-1,000
Buy Google	-2,000	2,100	-2,000
Don't invest	10	10	10

← **Maximin**

# Minimax

Select the minimum maximum regret

Action

State of Nature				Criterion
Poor market performance		Good market performance		Maximum regret for an action
Payoff	Regret	Payoff	Regret	
-1,000	1,010	1,700	400	1,010
-2,000	2,010	2,100	0	2,010
10	0	10	2,090	2,090

If I knew the future, which decision would I regret least?



# Minimax

Select the minimum maximum regret

Action

	State of Nature				Criterion
	Poor market performance		Good market performance		Maximum regret for an action
	Payoff	Regret	Payoff	Regret	
Buy Apple	-1,000	1,010	1,700	400	1,010
Buy Google	-2,000	2,010	2,100	0	2,010
Don't invest	10	0	10	2,090	2,090

←  
**Minimax**

If I knew the future, which decision would I regret least?

# Equal likelihood

Select the highest average payoff ASSUMING all states are of equal probability

Action	State of Nature		Criterion
	Poor market performance	Good market performance	Average reward/ payoff
	Payoff	Payoff	
Buy Apple	-1,000	1,700	350
Buy Google	-2,000	2,100	50
Don't invest	10	10	10
State Probability:	0.5	0.5	

# Equal likelihood

Select the highest average payoff ASSUMING all states are of equal probability

Action	State of Nature		Criterion
	Poor market performance	Good market performance	Average reward/ payoff
	Payoff	Payoff	
Buy Apple	-1,000	1,700	350
Buy Google	-2,000	2,100	50
Don't invest	10	10	10
State Probability:	0.5	0.5	

**Equal  
likelihood**



# Weighted average

Select the highest average payoff ASSUMING state probabilities from prior knowledge

Action	State of Nature		Criterion
	Poor market performance	Good market performance	Probability weighted average
	Payoff	Payoff	
Buy Apple	-1,000	1,700	80
Buy Google	-2,000	2,100	-360
Don't invest	10	10	10
State Probability:	0.6	0.4	

# Weighted average

Select the highest average payoff ASSUMING state probabilities from prior knowledge

	State of Nature		Criterion
	Poor market performance	Good market performance	Probability weighted average
	Payoff	Payoff	
Action	Buy Apple	-1,000      1,700	80
	Buy Google	-2,000      2,100	-360
	Don't invest	10      10	10
State Probability:	0.6	0.4	

**Weighted average**



# Decision making design pattern

1. Define a measure of risk or reward
2. Select the action that optimizes that metric

# Investments: reward

## State of Nature (s)

Action		Poor market performance	Excellent market performance
	Buy Apple	-1,000	1,700
	Buy Google	-2,000	2,100
	Don't invest	10	10

# Investments: reward

## State of Nature (s)

Action	State of Nature (s)	
	Poor market performance	Excellent market performance
	$s = s_0$	$s = s_1$
Buy Apple	-1,000	1,700
Buy Google	-2,000	2,100
Don't invest	10	10



# Investments: reward

## State of Nature (s)

		Poor market performance $s = s_0$	Excellent market performance $s = s_1$
Action	Buy Apple $a = a_0$	-1,000	1,700
	Buy Google $a = a_1$	-2,000	2,100
	Don't invest $a = a_2$	10	10

# Investments: reward

## State of Nature (s)

### Action

Poor market  
performance  
 $s = s_0$

Excellent market  
performance  
 $s = s_1$

Buy Apple  
 $a = a_0$

$V(a_0|s_0)$

-1,000

$V(a_0|s_1)$

1,700

Buy Google  
 $a = a_1$

$V(a_1|s_0)$

-2,000

$V(a_1|s_1)$

2,100

Don't invest  
 $a = a_2$

$V(a_2|s_0)$

10

$V(a_2|s_1)$

10

# Investments: reward

## State of Nature (s)

		Poor market performance $s = s_0$	Excellent market performance $s = s_1$
Action	Buy Apple $a = a_0$	$V(a_0 s_0)$ -1,000	$V(a_0 s_1)$ 1,700
	Buy Google $a = a_1$	$V(a_1 s_0)$ -2,000	$V(a_1 s_1)$ 2,100
	Don't invest $a = a_2$	$V(a_2 s_0)$ 10	$V(a_2 s_1)$ 10

**State Probability:**  $P(s_0) = 0.6$   $P(s_1) = 0.4$

# Investments: reward

$$EV(a_i) = V(a_i|s_0)P(s_0) + V(a_i|s_1)P(s_1)$$

↑  
Expected payoff

		State of Nature (s)	
		Poor market performance $s = s_0$	Excellent market performance $s = s_1$
Action	Buy Apple $a = a_0$	$V(a_0 s_0)$ -1,000	$V(a_0 s_1)$ 1,700
	Buy Google $a = a_1$	$V(a_1 s_0)$ -2,000	$V(a_1 s_1)$ 2,100
	Don't invest $a = a_2$	$V(a_2 s_0)$ 10	$V(a_2 s_1)$ 10

State Probability:  $P(s_0) = 0.6$        $P(s_1) = 0.4$

# Investments: reward

$$EV(a_i) = V(a_i|s_0)P(s_0) + V(a_i|s_1)P(s_1)$$

↑  
Expected payoff

		State of Nature (s)		Expected Reward $EV(a_i)$
		Poor market performance $s = s_0$	Excellent market performance $s = s_1$	
Action	Buy Apple $a = a_0$	$V(a_0 s_0)$ -1,000	$V(a_0 s_1)$ 1,700	
	Buy Google $a = a_1$	$V(a_1 s_0)$ -2,000	$V(a_1 s_1)$ 2,100	
	Don't invest $a = a_2$	$V(a_2 s_0)$ 10	$V(a_2 s_1)$ 10	
State Probability:		$P(s_0) = 0.6$	$P(s_1) = 0.4$	

# Investments: reward

$$EV(a_i) = V(a_i|s_0)P(s_0) + V(a_i|s_1)P(s_1)$$

↑  
Expected payoff

		State of Nature (s)		Expected Reward $EV(a_i)$
		Poor market performance $s = s_0$	Excellent market performance $s = s_1$	
Action	Buy Apple $a = a_0$	$V(a_0 s_0)$ -1,000	$V(a_0 s_1)$ 1,700	$(0.6)(-1000) + (0.4)(1700)$ <b>= 80</b>
	Buy Google $a = a_1$	$V(a_1 s_0)$ -2,000	$V(a_1 s_1)$ 2,100	
	Don't invest $a = a_2$	$V(a_2 s_0)$ 10	$V(a_2 s_1)$ 10	
State Probability:		$P(s_0) = 0.6$	$P(s_1) = 0.4$	

# Investments: reward

$$EV(a_i) = V(a_i|s_0)P(s_0) + V(a_i|s_1)P(s_1)$$

↑  
Expected payoff

		State of Nature (s)		Expected Reward $EV(a_i)$
		Poor market performance $s = s_0$	Excellent market performance $s = s_1$	
Action	Buy Apple $a = a_0$	$V(a_0 s_0)$ -1,000	$V(a_0 s_1)$ 1,700	$(0.6)(-1000) + (0.4)(1700)$ <b>= 80</b>
	Buy Google $a = a_1$	$V(a_1 s_0)$ -2,000	$V(a_1 s_1)$ 2,100	$(0.6)(-2000) + (0.4)(2100)$ <b>= -360</b>
	Don't invest $a = a_2$	$V(a_2 s_0)$ 10	$V(a_2 s_1)$ 10	
State Probability:		$P(s_0) = 0.6$	$P(s_1) = 0.4$	

# Investments: reward

$$EV(a_i) = V(a_i|s_0)P(s_0) + V(a_i|s_1)P(s_1)$$

↑  
Expected payoff

		State of Nature (s)		Expected Reward $EV(a_i)$
		Poor market performance $s = s_0$	Excellent market performance $s = s_1$	
Action	Buy Apple $a = a_0$	$V(a_0 s_0)$ -1,000	$V(a_0 s_1)$ 1,700	$(0.6)(-1000) + (0.4)(1700)$ <b>= 80</b>
	Buy Google $a = a_1$	$V(a_1 s_0)$ -2,000	$V(a_1 s_1)$ 2,100	$(0.6)(-2000) + (0.4)(2100)$ <b>= -360</b>
	Don't invest $a = a_2$	$V(a_2 s_0)$ 10	$V(a_2 s_1)$ 10	$(0.6)(10) + (0.4)(10)$ <b>= 10</b>
State Probability:		$P(s_0) = 0.6$	$P(s_1) = 0.4$	



# Risk = expected loss (cost)

**Loss:**  $\lambda(a_i | s_j) \triangleq$  Loss incurred by choosing action  $i$  and the state of nature being state  $j$

**Risk:** 
$$R(a_i) = \sum_{j=1}^{N_s} \lambda(a_i | s_j) P(s_j)$$

**Goal:** Select action  $i$  for which  $R(a_i)$  is minimum

# Investments: loss


$$\lambda(a_i) = \lambda(a_i|s_0)P(s_0) + V(a_i|s_1)(s_1)$$

↑  
Expected loss

		State of Nature (s)		Expected Loss $\lambda(a_i)$
		Poor market performance $s = s_0$	Excellent market performance $s = s_1$	
Action	Buy Apple $a = a_0$	$\lambda(a_0 s_0)$ 1,010	$\lambda(a_0 s_1)$ -1,690	$(0.6)(1010) + (0.4)(-1690)$ <b>= -70</b>
	Buy Google $a = a_1$	$\lambda(a_1 s_0)$ 2,010	$\lambda(a_1 s_1)$ -2,090	$(0.6)(2010) + (0.4)(-2090)$ <b>= 370</b>
	Don't invest $a = a_2$	$\lambda(a_2 s_0)$ 0	$\lambda(a_2 s_1)$ 0	$(0.6)(0) + (0.4)(0)$ <b>= 0</b>
State Probability:		$P(s_0) = 0.6$	$P(s_1) = 0.4$	


**How does this relate to supervised learning?**

# Where to operate along ROC?



# Where to operate along ROC?


		State of Nature	
		Class 0	Class 1
Estimate	Class 0	$\lambda_{00} = 0$	$\lambda_{01} = 100$
	Class 1	$\lambda_{10} = 1$	$\lambda_{11} = 0$



$\lambda_{ij}$  = Classify as class  $i$   
when state of nature is  
class  $j$

# Where to operate along ROC?


		State of Nature	
		Class 0	Class 1
Estimate	Class 0	$\lambda_{00} = 0$	$\lambda_{01} = 100$
	Class 1	$\lambda_{10} = 1$	$\lambda_{11} = 0$




$\lambda_{ij}$  = Classify as class  $i$   
when state of nature is  
class  $j$

- Assume our classification problem is landmine detection
- A false alarm wastes some time and resources, but a missed detection may cost a life

# Where to operate along ROC?



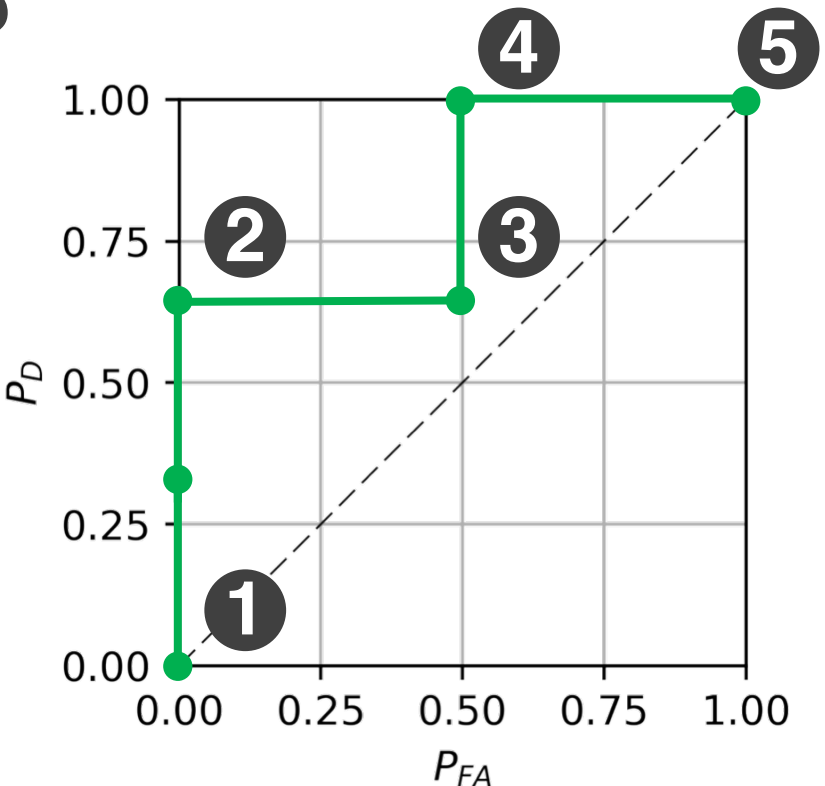
# Where to operate along ROC?



		State of Nature	
		Class 0	Class 1
Estimate	Class 0	$\lambda_{00} = 0$	$\lambda_{01} = 100$
	Class 1	$\lambda_{10} = 1$	$\lambda_{11} = 0$



# Where to operate along ROC?




$$R(a_i) = \lambda_{01}(1 - P_D(i)) + \lambda_{10}P_{FA}(i)$$

		State of Nature	
		Class 0	Class 1
Estimate	Class 0	$\lambda_{00} = 0$	$\lambda_{01} = 100$
	Class 1	$\lambda_{10} = 1$	$\lambda_{11} = 0$

# Where to operate along ROC?

Action: select operating point $i$	Probability of false alarm $P_{FA}$	Probability of missed detection $(1 - P_d)$	Risk $R(a_i)$
1	0	1	100
2	0	0.33	33
3	0.5	0.33	33.5
4	0.5	0	0.5
5	1	0	1



State of Nature

Class 0                      Class 1

$$R(a_i) = \lambda_{01}(1 - P_D(i)) + \lambda_{10}P_{FA}(i)$$


Estimate

Class 0  
  
Class 1

$\lambda_{00} = 0$	$\lambda_{01} = 100$
$\lambda_{10} = 1$	$\lambda_{11} = 0$

# Where to operate along ROC?

Action: select operating point $i$	Probability of false alarm $P_{FA}$	Probability of missed detection $(1 - P_d)$	Risk $R(a_i)$
1	0	1	100
2	0	0.33	33
3	0.5	0.33	33.5
4	0.5	0	0.5
5	1	0	1



State of Nature

Class 0                      Class 1

$$R(a_i) = \lambda_{01}(1 - P_D(i)) + \lambda_{10}P_{FA}(i)$$

Estimate

Class 0  
Class 1

$\lambda_{00} = 0$	$\lambda_{01} = 100$
$\lambda_{10} = 1$	$\lambda_{11} = 0$

**Can we use this theory to develop a classifier?**

# Binary decision

		State of Nature	
		Class 0 $s = s_0$	Class 1 $s = s_1$
Estimate	Class 0 $a = a_0$	$\lambda(a_0 s_0)$ $\lambda_{00}$	$\lambda(a_0 s_1)$ $\lambda_{01}$
	Class 1 $a = a_1$	$\lambda(a_1 s_0)$ $\lambda_{10}$	$\lambda(a_1 s_1)$ $\lambda_{11}$

$\lambda_{ij}$  = Loss when you classify as class  $i$  when state of nature is class  $j$

$$R(a_0|\mathbf{x}) = \lambda_{00}P(s_0|\mathbf{x}) + \lambda_{01}P(s_1|\mathbf{x})$$

$$R(a_1|\mathbf{x}) = \lambda_{10}P(s_0|\mathbf{x}) + \lambda_{11}P(s_1|\mathbf{x})$$

$$P(s_i|\mathbf{x}) = \frac{P(\mathbf{x}|s_i)P(s_i)}{P(\mathbf{x})}$$

**1**

Define the risk associated with each of the two actions

$$R(a_0|\mathbf{x}) = \lambda_{00}P(s_0|\mathbf{x}) + \lambda_{01}P(s_1|\mathbf{x})$$

$$R(a_1|\mathbf{x}) = \lambda_{10}P(s_0|\mathbf{x}) + \lambda_{11}P(s_1|\mathbf{x})$$

**2**

Create a decision rule based on the data

If  $R(a_0|\mathbf{x}) < R(a_1|\mathbf{x})$  then  $a_0$

Else  $R(a_0|\mathbf{x}) > R(a_1|\mathbf{x})$  then  $a_1$

**3**

Interpret that rule as a function of loss

$$\lambda_{00}P(s_0|\mathbf{x}) + \lambda_{01}P(s_1|\mathbf{x}) > \lambda_{10}P(s_0|\mathbf{x}) + \lambda_{11}P(s_1|\mathbf{x})$$

$$\frac{P(s_1|\mathbf{x})}{P(s_0|\mathbf{x})} > \frac{\lambda_{10} - \lambda_{00}}{\lambda_{01} - \lambda_{11}} \quad \text{then } a_1$$

# Minimizing the misclassification rate

$$\frac{P(s_1|\mathbf{x})}{P(s_0|\mathbf{x})} > \frac{\lambda_{10} - \lambda_{00}}{\lambda_{01} - \lambda_{11}} \quad \text{then } a_1$$

Assume that the loss is only for error, and it's the same for both types of error:

$$\lambda_{10} = \lambda_{01} \quad \text{and} \quad \lambda_{00} = \lambda_{11} = 0$$

Then the decision rule becomes:

$$\frac{P(s_1|\mathbf{x})}{P(s_0|\mathbf{x})} > 1 \quad \text{then } a_1$$

Pick whichever class is more likely given the data

# Recall Bayes Rule

$$\underset{\text{Posterior}}{P(s_i|\mathbf{x})} = \frac{\overset{\text{Likelihood}}{P(\mathbf{x}|s_i)} \overset{\text{Prior}}{P(s_i)}}{\underset{\text{Evidence}}{P(\mathbf{x})}}$$



**4**

Use Bayes rule to express this as a function of likelihoods

$$\frac{P(s_1|\mathbf{x})}{P(s_0|\mathbf{x})} > \frac{\lambda_{10} - \lambda_{00}}{\lambda_{01} - \lambda_{11}}$$

$$P(s_i|\mathbf{x}) = \frac{P(\mathbf{x}|s_i)P(s_i)}{P(\mathbf{x})}$$

$$\frac{P(\mathbf{x}|s_1)P(s_1)}{P(\mathbf{x}|s_0)P(s_0)} > \frac{\lambda_{10} - \lambda_{00}}{\lambda_{01} - \lambda_{11}}$$

then  $a_1$

**5**

The decision rule can be expressed as a **likelihood ratio**

$$\frac{P(\mathbf{x}|s_1)}{P(\mathbf{x}|s_0)} > \frac{\lambda_{10} - \lambda_{00}}{\lambda_{01} - \lambda_{11}} \frac{P(s_0)}{P(s_1)}$$

then  $a_1$

Class conditional  
likelihoods

False positive

$$\frac{P(\mathbf{x}|s_1)}{P(\mathbf{x}|s_0)} > \frac{\lambda_{10} - \lambda_{00} P(s_0)}{\lambda_{01} - \lambda_{11} P(s_1)}$$

then  $a_1$

Missed detection

**Estimate**

Class 0

Class 1

**State of Nature**

Class 0

Class 1

		State of Nature	
		Class 0	Class 1
Estimate	Class 0	$\lambda_{00} = 0$	$\lambda_{01} = 100$
	Class 1	$\lambda_{10} = 1$	$\lambda_{11} = 0$

# We can use this with generative models

$$\frac{P(\mathbf{x}|s_1)}{P(\mathbf{x}|s_0)} > \frac{\lambda_{10} - \lambda_{00}}{\lambda_{01} - \lambda_{11}} \frac{P(s_0)}{P(s_1)} \quad \text{then } a_1$$

- If we have class conditional models for our data, we can classify them
- **Naïve Bayes**: assume all features are independent given their class
- **Linear discriminant analysis**: assume the class conditional distributions, e.g.  $P(\mathbf{x}|s_1)$ , are each Gaussian

# Generative and discriminative models

Unobservable

**Data Generating Process**

$$p(X, Y)$$

**Types of models.** We can either model the full data generating process **OR** the target function, the mapping  $x$  to  $y$

→ If we model this process, it's a **generative model**

- Models  $P(x|y)$
- Can be used to generate synthetic data and impute missing values

**Target Function** for predicting  $y$  from  $x$

$$f(x) \rightarrow y$$

→ If we model this function, it's a **discriminative model**

- May model  $P(y|x)$  or directly map  $x$  to  $y$  without probabilities
- Often better performance for large sample sizes

# Takeaways

To make a decision:

1. Define a measure of risk or reward
2. Select the action that optimizes that metric

Decision theory informs how to operate supervised learning algorithms in practice

Decision theory incorporates relative importance of errors

Generative models estimate  $P(x|y)$ , while discriminative models estimate  $P(y|x)$